

Modern Big Data Analysis with SQL

Coursera Specialisation (Offered by Cloudera)

Course-3: Managing Big Data in Cluster and Cloud Storage

Week-1: Core

Question-1: Use the Table Browser in Hue to view the tables in the fun database. (See the environment installation instructions for how to log in.) Which of the following tables are in the fun database? Check all that apply.

Answer-1: card_rank

- card_suit
- games
- inventory

Question-2: Use the Table Browser in Hue to view the columns in the employees table in the default database. Which of the following columns are in the employees table? Check all that apply.

Answer-2: empl_id

- first_name
- last_name
- office_id
- salary

Question-3: Which of the following is a way to show what tables are available in the fly database using the query editor (center panel of Hue)?

Answer-3: Select the fly database as the active database and then run the command SHOW TABLES;

Question-4: By default, what directory in HDFS would store the data for a table named artists in a database named music?

Answer-4: /user/hive/warehouse/music.db/artists

Question-5: By default, what directory in HDFS would store the data for a table named facilities in the default database?

Answer-5: /user/hive/warehouse/facilities

Question-6: What delimiter is used to separate the values in the lines of the text file containing the data in the orders table in the default database?

Answer-6: tab (\t)

Question-7: Which command will list the files in an HDFS directory to the terminal screen?

Answer-7: hdfs dfs -ls /path/to/directory/

Question-8: Which is an advantage of using S3?

Answer-8: Scalability is easier with S3 (or other cloud storage) than on-premises storage

Question-9: Which commands can you use to print the contents of a file on S3 to your terminal screen? Check all that apply.

Answer-9: `hdfs dfs -cat`