

Modern Big Data Analysis with SQL

Coursera Specialisation (Offered by Cloudera)

Course-3: Managing Big Data in Cluster and Cloud Storage

Week-4: Core

Question-1: Which of these are ways to load data into HDFS using Hue on the VM?

Check all that apply.

Answer-1: Load data into any HDFS directory through data source assist panel on left side

- Use the File Browser to load data into any HDFS directory
- Use the Table Browser to load data into the storage directory of an existing table.
- Use the Table Browser to load data into the storage directory of a new table at a time when you create the table

Question-2: Which of these tasks can you complete using HDFS shell commands? Check all that apply.

Answer-2: Load data into the storage directory of an existing table

- Load data for a storage directory for a table that hasn't been created yet

Question-3: Which hdfs dfs command option can be used to upload a file from your local filesystem to HDFS?

Answer-3: -put

Question-4: On the course VM, which command could you use to upload the local file /home/training/training_materials/analyst/data/games.csv to an S3 bucket named bucket-o-games? For purposes of this question, assume you have write access to this bucket.

Answer-4: `aws s3 cp /home/training/training_materials/analyst/data/games.csv s3:// bucket-o-games`

Question-5: What is the effect of including TBLPROPERTIES ('serialization.null.format' = '') in a CREATE TABLE statement?

Answer-5: Empty strings in the table's data files will appear as NULL in query results

Question-6: This command uses Sqoop to import data from MySQL to HDFS on the VM (with the user training).

```
$ sqoop import --connect jdbc:mysql://localhost/mydb \  
--username training --password training \  
--table example_table \  
--target-dir /user/hive/warehouse/example_table
```

Where in HDFS will the data be found?

Answer-6: /user/hive/warehouse/example_table

Question-7: Which are valid uses for Sqoop? Check all that apply.

Answer-7: Exporting data from HDFS to an RDBMS (such as MySQL)

- Importing data to HDFS from an RDBMS (such as MySQL)

Question-8: Which are advantages of using the LOAD DATA SQL statement instead of the `hdfs dfs -mv` shell command to load data? Check all that apply.

Answer-8: LOAD DATA renames any files that are the same as files that already exist within the directory.

- If you use Impala to execute the LOAD DATA command, the metadata cache is automatically updated

Question-9: Why is it a poor practice to use INSERT commands to load data into a table a few rows at a time?

Answer-9: You get many small files, which can slow big data systems down

Question-10: Which are valid ways to create a copy of an existing table? Check all that apply.

Answer-10: `CREATE TABLE newtable LIKE oldtable;` then `* INSERT INTO newtable`

`SELECT * FROM OLDTABLE;`

`CREATE TABLE newtable AS SELECT * FROM oldtable;`